

Run book for SolrCloud configuration in K8S.

Step 2: Make sure your K8S running

Step 2: Install SolrCloud in k3s.

```
git clone https://github.com/freedev/solrcloud-zookeeper-kubernetes.git
```

```
cd solrcloud-zookeeper-kubernetes
```

Step 3: Make below changes to add 2 nodes to single solrcloud cluster

- 1) Navigate to statefulsets folder and modify below files to change replica to 3.

Statefulset-zookeeper-ensemble.yml

Statefulset-zookeeper.yml

Statefulset-solr-cluster.yml (change to 4)

Statefulset-solr.yml

- 2) Navigate to configmap folder and modify solr-config.properties to add other zookeeper details in ensemble

zkHost=zk-0.zkensemble:2181, zk-1.zkensemble:2181, zk-2.zkensemble:2181

- 3) Navigate to minikube folder and change the app version to V1 in below folders

Storageclass-solrcluster.yml

Storageclass-zkensemble.yml

- 4) Run below command to create/deploy SolrCloud cluster

```
./start-minikube.sh
```

Run below query to check pods

```
Kubectl get pods
```

You can access solar master at **http://< external ip>:8983/solr**

or

Run minikube service list and click on the URL for SolrCloud.

Once SolrCloud is up and running,

To change the solr configuration, we need to login into the pod.

```
Kubectl -it exec podname -c containername bash
```

Eg: kubectl -it exec solr-0 -c solr bash.

Collection creation

- 1) Create a Collection in API using 2 shards and 2 replicas.

My_collection

- 2) Use maxshards per node as 2

The screenshot shows the MongoDB Compass interface for a collection named 'my_collection'. The configuration is as follows:

- Shard count: 2
- configName: _default
- replicationFactor: 2
- maxShardsPerNode: 2
- router: compositeltd
- autoAddReplicas: true

On the right, the shards are listed as 'Shard: shard1' and 'Shard: shard2'.

- 3) Once collection is create we can load using post

Post -c my_collection data files

Eg : post -c mycollection example/exampledocs/*

- 4) Once created you can check the data in queries

The screenshot shows the MongoDB Compass interface with a query result for a collection named 'my_collection'. The query is 'q: {}' and the result is a JSON array of documents:

```
{
  "id": "0812521390",
  "cat": "book",
  "name": "The Black Company",
  "price": 6.99,
  "inStock": false,
  "author": "Glen Cook",
  "series_t": "The Chronicles of The Black Company",
  "sequence_i": 1,
  "genre_s": "fantasy",
  "_version_": 17103503796870464,
}
{
  "id": "0441385532",
  "cat": "book",
  "name": "Jhereg",
  "price": 7.95,
  "inStock": false,
  "author": "Steven Brust",
  "series_t": "Vlad Taltos",
  "sequence_i": 1,
  "genre_s": "fantasy",
  "_version_": 1710350380012929024,
}
{
  "id": "0380014300",
  "cat": "book",
  "name": "Nine Princes In Amber",
}
```

- 5) You can the shards and replicas and how many documents are loaded into each using nodes in Admin.

Host	Node	CPU	Heap	Disk usage	Requests	Collections	Replicas
solr-0.solrcluster Linux 5.4Gb java 11 Load: 1.43 show details...	8983_solr Uptime: 16m show details...	0%	31%	48.3Kb	RPM: 0.04 p95: 912ms	my_collection	my_collection_s2r4 (29 docs)
solr-1.solrcluster Linux 5.4Gb java 11 Load: 1.43 show details...	8983_solr Uptime: 15m show details...	0%	66%	50.1Kb	RPM: 0.05 p95: 310ms	my_collection	my_collection_s1r3 (23 docs)
solr-2.solrcluster Linux 5.4Gb java 11 Load: 1.43 show details...	8983_solr Uptime: 15m show details...	0%	57%	96.1Kb	RPM: 0.02 p95: 364ms	my_collection	my_collection_s1r1 (23 docs) my_collection_s2r6 (29 docs)